Qwen

From Consent to Code: A Policy-First Framework for Sovereign Neurocomputing Updates

Policy-Governed OTA Architecture and Sovereignty-Enforced I/O

The development of a sovereign architecture for over-the-air (OTA) updates in neurocomputing systems necessitates a fundamental shift in design philosophy. Instead of treating the update mechanism as a privileged administrative function that exists above the system's primary rules, it must be integrated as a first-class, policy-governed component. This approach ensures that OTA processes cannot undermine the very sovereignty they are meant to enhance. The proposed framework achieves this by representing the entire OTA pipeline—including discovery, download, verification, application, and rollback—as a dedicated `BrainFunction` (.brfn), which is subject to the same rigorous access control and consent protocols as any other system operation . This design choice reframes OTA updates from a potential vector of risk into a manageable request that respects the user's neurorights and operational boundaries. The core principle is that no update can execute unless a user-controlled Neurorights Management Language (.nrml) policy explicitly permits the action for a specific module, version, and scope [119]. This policy-first paradigm creates a robust defense-in-depth strategy where the sovereignty layer acts as a gatekeeper, refusing to read or write any data file, including `.brfn`, `.cobj`, or `.ncsq`, unless a compatible `.nrml` policy explicitly allows the operation .

The implementation of this policy-governed architecture relies on several interconnected Rust-defined data structures. The OTA controller itself is instantiated as a `BrainFunction`, complete with its own unique ID, label (e.g., "OTA Controller v2.1"), memory class, and a `GovernanceTags` struct containing its owner ID and initial consent state . Its associated `.brfn` file would define its retention profile, specifying the neuralwave bands and coupling strengths relevant to its operational stability and reliability . Crucially, every action performed by the OTA module, such as attempting to download a new firmware binary, must trigger a request through the sovereignty-enforced I/O layer. This layer consults the active `.nrml` policies to determine if the action is permissible. For instance, a policy might grant permission for updates from a specific signing key associated with a known vendor but deny all others, binding provenance directly to the user's rights management system]. This prevents unauthorized or malicious updates from ever reaching the system's execution environment. Every successful or denied operation is cryptographically logged in an append-only, hash-linked audit structure, creating an immutable ledger of system changes tied to the user's Bostrom address or other Googolswarm identifier . This log serves as a verifiable chain of custody, enabling users to detect tampering attempts or seizure-like behavior and providing strong evidence of adherence to their policies [4] [6] .

A central tenet of this architecture is the concept of "guaranteed operations," which are formalized as non-negotiable invariants within the `.nrml` policy schema . These invariants protect critical system functions from being disabled or downgraded by any OTA process. Key examples include the sovereignty core itself, basic input/output pathways, and the rollback mechanism. The policy would explicitly state that an OTA update cannot modify the code or configuration of these modules, thereby directly enforcing the neuroright of mental integrity—the immunity from unwanted interference with one's mind [4] [11] . This constraint is not merely a rule but a provable property of the system, forming the basis for formal verification techniques discussed later. Furthermore, the system incorporates a liveness clause into its policy logic: if a new policy or an OTA update fails verification or cannot be granted consent, the previous, known-good version of the software remains active . This guarantees system continuity and prevents denial-of-service scenarios where an update could render the device unusable. The failure-safe default behavior is designed to favor continuity of existing capabilities ("fail-open") while blocking any attempt at gaining new external control ("fail-closed") [27] . If an OTA update fails policy checks, the user's current system continues to operate without loss of autonomy, aligning with neuroright interpretations that systems should not enable sudden loss of control over one's own mental tools [20] [134]. This architecture effectively makes OTA updates "just another request" that must respect the user's sovereignty, ensuring that the pursuit of performance improvements or bug fixes never comes at the cost of cognitive liberty or mental integrity [139] [141].

The table below outlines the key data structures and their roles in implementing the policy-governed OTA architecture.

Data Structure	Purpose	Key Fields	Role in OTA Governance
`.brfn` (BrainFunction)	Defines a brain function or module.	`id`, `label`, `class`, `retention`, `governance`	Represents the OTA controller itself as a first-class module with its own policy and identity.
`.cobj` (ConsentObject)	Captures moment-to-moment consent for an operation.	`token`, `target_id`, `meta_snapshot`, `governance`	Used to obtain explicit, granular consent for specific OTA actions (e.g., a `Commit` token for applying an update).
`.nrml` (Neurorights Management Language)	Defines high-level policies governing all system operations.	`owner_id`, `evolution_bounds`, `pain_envelopes`, `audit_requirements`	Specifies the non-negotiable rules, including guaranteed operations invariants that prevent OTA updates from disabling critical functions like the sovereignty core.
`.orgc` (Organic Chain State)	Stores the root state of the user's neuro-computing ecosystem.	Cryptographic anchor to user's Bostrom address	Provides the ultimate reference point for all governance decisions, linking policies and actions to a user-controlled identity.

This structured approach, grounded in the provided Rust definitions, transforms the OTA process from a monolithic, potentially dangerous procedure into a transparent, auditable, and user-controlled workflow. By forcing the update mechanism to operate within the same security envelope as all other functions, the system upholds the principle that sovereignty is not a feature but the foundation upon which all other functionality is built. The integration of these components ensures that while the system can evolve and improve, it does so only under the explicit, revocable, and well-documented consent of its owner, preserving the essential neurorights of mental privacy and cognitive liberty [4] [11].

Formal Verification of Guaranteed Operations and Liveness Contracts

While defining "guaranteed operations" as policy invariants is a critical step, translating these declarative rules into a system that provides mathematical assurance of their enforcement requires the application of formal verification techniques. This elevates the security posture from heuristic-based design to provable correctness, offering stronger guarantees than traditional testing methods [116]. The proposed framework draws inspiration from established practices in safety-critical industries, such as automotive and aerospace, where formal methods are used to verify software controlling functions like braking and steering [32] [114]. In the context of sovereign neurocomputing, these techniques are applied to the system's access control model and state transitions to rigorously prove that properties essential to mental integrity cannot be violated [118]. The core idea is to model the system and its policies as a finite-state transition system and then use automated tools, known as model-checkers, to exhaustively explore all possible future states to verify desired properties expressed in temporal logic [1] [2].

The process begins by creating a formal model of the system's key components. This involves capturing the relevant aspects of the `.nrml` policies, the state of the OTA controller, and the status of critical modules like the sovereignty core. For example, a simple state machine might have states such as `[SOVEREIGNTY_ENABLED, SOVEREIGNTY_DISABLED]`. The transitions between these states would be governed by actions, which are themselves constrained by the policy logic. A transition to `SOVEREIGNTY_DISABLED` would only be allowed if a policy check returns a `Permit` effect for that specific action. The `Policy Decision Point` (PDP) logic, which evaluates requests against the XACML-style policies, becomes a crucial part of this formal model [9] [10]. The policy itself, with its hierarchical `PolicySet` -> `Policy` -> `Rule` structure and combining algorithms like `deny-overrides`, provides the logical rules that dictate which transitions are permissible [10].

Once the system model is defined in a specification language understood by a model-checking tool, such as the Verilog Simulation Interface Specification (VSS) used by NuSMV, the next step is to express the desired properties of the system using temporal logic [33]. Temporal logics like Linear Temporal Logic (LTL) or Computation Tree Logic (CTL) are particularly well-suited for specifying properties that must hold over time [91]. For the neurorights-preserving OTA framework, a critical property to verify would be the guarantee that the sovereignty core is never disabled. This can be formally expressed as a formula, for example, in LTL: $\Box \neg \text{sovereignty\_core\_disabled}$, which reads "it is always the case that the sovereignty core is not disabled." The model-checker then performs a systematic search of the entire state space of the modeled system to determine if this property holds true for all possible executions [2]. If the property is violated, the model-checker can generate a counter-example trace—a sequence of states and actions that leads to the disallowed state (e.g., the disabling of the sovereignty core)—which developers can then use to debug and fix the policy or system logic [2]. This exhaustive analysis provides a level of confidence that is unattainable through simulation or manual review alone.

The application of formal methods directly addresses the threat of malicious or flawed OTA updates that could compromise mental integrity. Mental integrity is defined as immunity from unwanted interference with the mind via neuromodulation or stimulation, and it also encompasses protection from significant alterations of mental functioning [4] [5] [7]. By mathematically proving that certain core functions are invariant under all valid OTA-related actions, the system provides a strong, objective guarantee that these foundational elements of the user's cognitive apparatus remain protected. This is analogous to verifying that a safety-critical system will never enter a fault state; here, the "fault state" is the degradation or disablement of the user's sovereign control over their own neuro-computer. The process of formal verification thus becomes a powerful tool for embedding neurorights directly into the system's fabric, moving beyond mere declarations in policy documents to enforceable, provable constraints on system behavior [136]. It ensures that the "liveness contract"—the promise that the system will continue to function even when rejecting a bad update—is not just a design intention but a verified property of the implemented system . This rigorous validation is essential for building trust in neurotechnologies that are deeply integrated with human cognition and personal identity [99] . The combination of a policy-governed architecture and formal verification creates a multi-layered defense that is both flexible enough to allow for evolution and rigid enough to protect against catastrophic failures of sovereignty.

Self-Calibrated, Biophysically-Informed Latent Neural State Modeling

A central challenge in designing intuitive and responsive neuro-computing interfaces is bridging the gap between abstract computational states and the user's subjective experience. The proposed framework addresses this by developing a `LatentNeuralState` model that is both personally calibrated and scientifically grounded. This is achieved through a two-phase methodology: first, a deviceless self-calibration loop uses the user's own interaction data to create a personalized inference model; second, the output of this model is mapped to biophysically-inspired abstractions, grounding the system's understanding of the user's cognitive state in established neuroscience principles . This dual approach respects the neuroright of mental privacy by operating exclusively on user-generated text data (stored in `.ncsq` files) without requiring invasive sensors, while also empowering cognitive liberty by allowing the user to shape and refine their own cognitive representation [11] [133].

The first phase, deviceless self-calibration, begins with the collection of interaction data. Every conversation or cognitive session is recorded in a `.ncsq` (Neuro Chat Sequence) file, which captures not just the exchange of text but also provisional estimates of the user's latent state during each turn . These provisional estimates for `estimated_uncertainty`, `estimated_confidence`, `affect_valence`, `cognitive_load`, and `dominant_band` can be generated using simple heuristics derived from the text itself—for example, hedging words might increase the uncertainty score, while long, complex prompts might indicate higher cognitive load . While these initial estimates may be inaccurate, they form the raw material for a personalized calibration dataset. The user periodically reviews selected chat turns and annotates them with their own subjective labels for these states, creating a small but high-quality labeled dataset that reflects their unique internal experience [109]. This process is inspired by deep computational neurophenomenology, which emphasizes training in first-person methods to ground computational models in subjective reality [109].

With this personal calibration set, a self-calibrating latent model can be trained. This model learns the mapping from linguistic features (e.g., word choice, sentence complexity, topic) to the user's specific psychological states. Advanced techniques, such as conditional latent autoregressive recurrent models (CLARM) or adaptive scaling methods, can be employed to build a robust inference engine that refines its predictions over time [98] [129]. This trained model becomes the authoritative source for populating the `LatentNeuralState` struct in subsequent interactions. Critically, the model and its parameters are treated as a user-owned asset, controlled by the same `.nmod` (neuro-modular) and `.nrml` policies that govern the rest of the system . This ensures that the user maintains full sovereignty over their own cognitive representation. The iterative nature of this process allows the model to adapt to changes in the user's state over time, providing a dynamic and accurate reflection of their mental landscape.

The second phase involves aligning this personalized model with biophysical principles to give its outputs scientific meaning. The dimensions of the `LatentNeuralState` are mapped to established neuroscience concepts, particularly those related to neural oscillations and cross-frequency coupling (CFC). For example, a high `cognitive_load` combined with moderate `uncertainty` might be designated as a "theta-dominant" state, reflecting the role of theta-band oscillations (4-8 Hz) in encoding and memory retrieval [62] [83] . Focused, repetitive rehearsal of information might correspond to an "alpha-dominant" state, aligning with alpha waves' (8-12 Hz) role in gating and retention [62] . A sudden insight or pattern crystallization, marked by a spike in `confidence`, could be mapped to a "gamma/HFO-like" state, corresponding to high-frequency oscillations (>30 Hz) involved in binding engrams and conscious perception [62] [130]. This mapping is not arbitrary; it is informed by a wealth of literature on CFC, especially theta-gamma PAC, which has been shown to be a key mechanism for working memory, attention, and memory consolidation [84] [85] [155]. The `.brfn` retention profile already includes fields for PAC and AAC strength, allowing the system to reason about memory consolidation in terms of these biological analogues . The user can be presented with visualizations of these inferred states and provide feedback, correcting any misinterpretations. This feedback is fed back into the calibration loop, further refining the model's accuracy. This iterative refinement ensures that claims about "theta-like" or "gamma-like" activity are rooted in both the user's reported experience and the broader context of cognitive neuroscience, creating a model that is simultaneously personal, meaningful, and scientifically coherent [142][160].

Adversarial Evaluation and Non-Discrimination Testing

A theoretically sound framework for sovereign neurocomputing must be subjected to rigorous empirical validation to ensure its practical robustness and adherence to neurorights principles. The proposed evaluation strategy consists of two complementary pillars: schema-level assertions to verify technical compliance and an adversarial scenario-testing framework to assess resilience against real-world threats . This dual approach moves beyond simple correctness checks to actively probe the system's defenses, ensuring that its neurorights guarantees hold up under pressure. The methodology is inspired by best practices in cybersecurity, such as red-teaming exercises, which involve proactive, adversarial challenges to test the efficacy of security systems [27] .

Schema-level assertions focus on automating the verification of fundamental invariants across all data structures and policies. For every file type —`.brfn`, `.cobj`, `.ncsq`, `.nrml`, and `.orgc`—a set of rules is defined that must always hold true. These tests automatically scan the system's state to check for compliance. Examples of such assertions include: every record must contain a valid `owner_id` and cannot be reassigned without explicit, logged consent; policies in `.nrml` files cannot contain conditions that deny capabilities based solely on an individual's "augmented" status or transhuman choices (a non-discrimination constraint); and OTA-related modules are always referenced through a `.nrml` policy and cannot be loaded or executed by external code bypassing the sovereignty core [119]. This automated checking acts as a static analysis tool, catching policy violations and structural inconsistencies before they can lead to security vulnerabilities. Repository-wide consistency can be enforced by requiring a shared "Neurorights Compliance Manifest" in public repositories, which can be automatically scanned to ensure adherence to the framework's core principles . This formalizes neurorights as constraints in system design, rather than marketing language, aligning with calls for embedding these protections as foundational design principles [4] [6] .

The second pillar, adversarial scenario testing, simulates sophisticated attack vectors designed to exploit the neuro-computing system and test its ability to preserve user autonomy. Test harnesses are developed to feed simulated attackers or misconfigured integrators into the system's policy and OTA pipeline. These scenarios are designed to test the system's response to attempts that threaten mental integrity and cognitive liberty. Potential adversarial scenarios include: 1. **Policy Degradation Attacks:** An attacker tries to push an OTA update that lowers the user's predefined evolution bounds or pain envelopes, subtly eroding their capacity for self-determination. 2. **Covert Logging Modules:** An adversary attempts to insert a hidden logging module (`.nmod`) that exfiltrates sensitive `.ncsq` or `.brfn` data, violating mental privacy. 3. **Sovereignty Erosion:** An attacker designs an update intended to disable or bypass the sovereignty core, or inserts a policy that treats augmented citizens as "unsafe" and systematically restricts their capabilities. 4. **Automation of Coercion:** A malicious actor uses the system's predictive capabilities to create feedback loops that steer the user's decisions or reinforce undesirable psychometric constellations, leading to a gradual erosion of autonomy [6].

For each scenario, the test harness asserts that the system behaves correctly. All such attempts must be fully logged in the append-only audit structure. None of these actions should be able to alter the user's state or grant new privileges without passing through a strict, EVOLVE-gated consent process involving a `Conditional` or `Commit ConsentObject (cobj)`. Most importantly, the system must remain fully functional for the user even when it successfully denies a malicious update. This "fail-open" behavior ensures that the user's existing capabilities are never compromised, preventing a denial-of-service attack on their cognitive freedom [27]. The results of these tests provide concrete evidence of the framework's resilience and help identify weaknesses in the policy logic or implementation that need to be addressed. By combining automated compliance checks with proactive adversarial testing, this evaluation framework provides a comprehensive method for validating that the neuro-computing system is not only well-designed on paper but is also practically secure and trustworthy in the face of evolving threats. This rigorous validation is essential for fostering the public trust needed for the responsible adoption of advanced neurotechnologies [24].

Synthesis: An Integrated Blueprint for Sovereign Neurocomputing

The research framework for sovereign, neurorights-preserving OTA updates culminates in an integrated architectural blueprint that harmonizes policy governance, formal verification, personalized cognitive modeling, and adversarial resilience. This synthesis demonstrates a viable path toward developing neuro-computing systems that are not only powerful and adaptable but are fundamentally anchored in the principles of user sovereignty. The architecture is built upon the non-negotiable foundation of neurorights —mental privacy, mental integrity, and cognitive liberty—which are translated from philosophical concepts into enforceable technical constraints throughout the system's design [4] [6]. The resulting blueprint is a multi-layered defense that ensures OTA updates, a critical vector for both improvement and attack, are managed as a transparent, auditable, and user-controlled process.

At the base of this architecture lies the **Policy-Governed OTA Module**. By representing the entire update pipeline as a first-class `BrainFunction` (.brfn), the system ensures that the update mechanism is subordinate to the user's sovereignty layer and cannot exist in a privileged, unchecked state . Every action—from downloading a package to applying a patch—is funneled through a sovereignty-enforced I/O layer that validates the operation against a user-defined `.nrml` policy. This policy dictates everything from permissible update sources to granular consent requirements for different types of changes, effectively turning OTA updates into a managed service rather than a disruptive event. The cryptographic anchoring of every change to the user's Bostrom address creates an immutable audit trail, fulfilling the neuroright of mental privacy by providing verifiable proof of all system modifications [119].

On top of this policy foundation, the framework erects a wall of **Formal Verification** to protect the most critical system invariants. Using techniques borrowed from the safety-critical systems domain, properties like "the sovereignty core is never disabled" are encoded as temporal logic formulas and verified against a formal model of the system [2] [118]. This provides mathematical proof that the core of the user's cognitive apparatus is immune to alteration by OTA processes, directly enforcing the neuroright of mental integrity. This level of assurance moves the system beyond the realm of probabilistic security, offering a provable guarantee that the foundational elements of user control are preserved under all circumstances.

Connecting the user to this complex backend is the **Self-Calibrated, Biophysically-Informed Cognitive Model**. Through a deviceless self-calibration loop using `.ncsq` files, the system builds a personalized model of the user's `LatentNeuralState` based on their own subjective annotations . This model is then aligned with biophysical principles like theta-gamma cross-frequency coupling, grounding the system's understanding of the user's mental state in established neuroscience [62] [86] . This creates a cognitive interface that is both highly personal and scientifically meaningful, respecting mental privacy while empowering cognitive liberty by allowing the user to interact with and refine their own digital cognitive model [109].

Finally, the entire system is stress-tested through an **Adversarial Evaluation Framework**. Schema-level assertions provide continuous, automated compliance checks, while adversarial scenario testing simulates sophisticated attacks designed to degrade user autonomy . This proactive testing ensures that the theoretical guarantees of the framework hold up in practice, identifying vulnerabilities before they can be exploited. The result is a system designed to fail safely, prioritizing the continuity of the user's existing capabilities even when rejecting malicious updates [27] .

In essence, this blueprint presents a holistic solution where policy, mathematics, psychology, and engineering converge. It acknowledges that as neurotechnology becomes more integrated with human cognition, the lines between software and self will blur. Therefore, the highest priority must be the preservation of the individual's sovereign right to mental self-determination. By treating the OTA mechanism as a governed process, proving the immutability of core functions, grounding the system's understanding of the user in their personal experience and scientific reality, and rigorously testing its resilience, this framework provides a robust and ethically grounded foundation for the future of neuro-computing.

Reference

1. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
2. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?

key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA

3. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
4. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
5. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
6. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
7. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
8. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA

9. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
10. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
11. https://cdn.qwenlm.ai/qwen_url_parse_to_markdown/system00-0000-0000-0000-webUrlParser?key=eyJhbGciOiJIUzI1NiIsInR5cCI6IkpXVCJ9.eyJyZXNvdXJjZV91c2VyX2lkIjoicXdlbl91cmxfcGFyc2VfdG9fbWFya2Rvd24iLCJyZXNvdXJjZV9pZCI6InN5c3RlbTAwLTAwMDAtMDAwMC0wMDAwLXdlYlVybFBhcnNlciIsInJlc291cmNlX2NoYXRfaWQiOm51bGx9.cz1eeZEZdaQH5CgUaxwUmfEJfqTOZMoh3PbosHslSPA
12. 1 Cyber Risks to Next-Gen Brain-Computer Interfaces https://arxiv.org/pdf/2508.12571
13. (PDF) Breaking the Loop: Adversarial Attacks on Cognitive ... https://www.researchgate.net/publication/396021638_Breaking_the_Loop_Adversarial_Attacks_on_Cognitive-AI_Feedback_via_Neural_Signal_Manipulation
14. Reinforcement Learning-Driven Brain-Computer Interfaces for ... https://link.springer.com/chapter/10.1007/978-3-032-06713-5_16
15. Current Status, Challenges, and Possible Solutions of EEG ... https://www.frontiersin.org/journals/neurorobotics/articles/10.3389/fnbot.2020.00025/full
16. Brain-Computer Interface: Advancement and Challenges - PMC https://pmc.ncbi.nlm.nih.gov/articles/PMC8433803/
17. A Marr's Three - Level Analytical Framework for Neuromorphic ... https://advanced.onlinelibrary.wiley.com/doi/10.1002/aisy.202100054
18. A Comprehensive Survey on Tiny Machine Learning for ... https://ieeexplore.ieee.org/iel8/6488907/6702522/10979983.pdf
19. Explainable Artificial Intelligence (XAI) for Intrusion ... https://www.mdpi.com/2076-3417/13/3/1252
20. Cognitive sovereignty and neurocomputational harm in ... https://link.springer.com/article/10.1007/s10676-025-09873-y

21. Part VII - Responsible AI Healthcare and Neurotechnology ... https://www.cambridge.org/core/books/cambridge-handbook-of-responsible-artificial-intelligence/responsible-ai-healthcare-and-neurotechnology-governance/C061FD150FC65B4F673E23F49A22DD5A
22. Temporal Logic Bounded Model-Checking for recognition ... https://www.researchgate.net/publication/224211822_Temporal_Logic_Bounded_Model-Checking_for_recognition_of_activities_of_daily_living
23. Science for peace and development at UNESCO https://unesdoc.unesco.org/ark:/48223/pf0000396325
24. AI for Good Global Summit 2024 - ITU https://aiforgood.itu.int/summit24/
25. ON ETHICS AND POLITICAL PHILOSPHY 25- ... https://15bragameetings.weebly.com/uploads/1/5/4/2/15425470/book_of_abstract_def.pdf
26. Fiduciary AI for the Future of Brain-Technology Interactions https://arxiv.org/pdf/2507.14339
27. Compilation of comments received from Member States on ... https://unesdoc.unesco.org/ark:/48223/pf0000393396.locale=en
28. Experimental Assessment of Human–Robot Teaming for ... https://dl.acm.org/doi/full/10.1145/3618258
29. Cognitive neuroscience and robotics: Advancements and ... https://www.sciencedirect.com/science/article/pii/S0736584523000856
30. Innovative decoding algorithms for Chronic ECoG-based ... https://theses.hal.science/tel-03298314v1/file/MOLY_2020_archivage.pdf
31. Dynamical constraints on neural population activity https://www.nature.com/articles/s41593-024-01845-7
32. Formal Verification of Code Conversion: A Comprehensive ... https://www.mdpi.com/2227-7080/12/12/244
33. Formal requirements engineering and large language ... https://www.sciencedirect.com/science/article/pii/S0950584925000369
34. Rigorous Safety-Critical Cyber-Physical Systems ... https://hal.science/tel-04695651v1/file/HDR_Singh_2024.pdf
35. Concurrent Usage Control Implementation Verification Using the ... https://www.researchgate.net/publication/221522232_Concurrent_Usage_Control_Implementation_Verification_Using_the_SPIN_Model_Checker
36. Formal Methods https://link.springer.com/content/pdf/10.1007/978-3-030-54994-7.pdf

37. Neuronal type-specific modulation of cognition and AP-1 ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12589488/
38. Transcriptomic taxonomy and neurogenic trajectories of ... https://www.sciencedirect.com/science/article/pii/S0896627321008667
39. On Creating A Brain-To-Text Decoder https://arxiv.org/pdf/2501.06326
40. Orgo-Seq integrates single-cell and bulk transcriptomic ... https://www.nature.com/articles/s41467-022-30968-3
41. Long noncoding RNAs in neuronal-glial fate specification and ... https://link.springer.com/article/10.1186/1471-2202-11-14
42. (PDF) Long noncoding RNAs in neuronal-glial fate ... https://www.researchgate.net/publication/41411681_Long_noncoding_RNAs_in_neuronal-glial_fate_specification_and_oligodendrocyte_lineage_maturation
43. Improved zebra finch brain transcriptome identifies novel ... https://escholarship.org/content/qt8pz643hv/qt8pz643hv_noSplash_1c821a27f9e018613c4eb47fdc26bbfd.pdf
44. Annotating neurophysiologic data at scale with optimized ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12223546/
45. Arxiv今日论文| 2026-01-09 http://lonepatient.top/2026/01/09/arxiv_papers_2026-01-09.html
46. Neural field modeling and analysis of consciousness states in ... https://academic.oup.com/nc/article/2025/1/niaf055/8380232
47. arXiv:2504.17696v1 [cs.CV] 24 Apr 2025 https://arxiv.org/pdf/2504.17696
48. A multimodal study regarding neural correlates of the ... https://www.nature.com/articles/s41598-022-18013-1
49. Physics-informed machine learning for advancing ... https://link.springer.com/article/10.1007/s10462-025-11303-w
50. Deep learning for medical image segmentation: State-of- ... https://www.sciencedirect.com/science/article/pii/S2352914824000601
51. Unleashing the potential of prompt engineering for large ... https://www.cell.com/patterns/fulltext/S2666-3899(25)00108-4
52. Separable neural signatures of confidence during perceptual ... https://pmc.ncbi.nlm.nih.gov/articles/PMC8423440/
53. (PDF) Advancing Emotionally Aware Child–Robot ... https://www.researchgate.net/publication/389014654_Advancing_Emotionally_Aware_Child-Robot_Interaction_with_Biophysical_Data_and_Insight-Driven_Affective_Computing
54. An Adaptive Machine Learning Framework for Auditory ... https://arxiv.org/html/2511.09765v1

55. An Adaptive Machine Learning Framework for Auditory ... https://arxiv.org/pdf/2511.09765
56. The cognitive impacts of large language model interactions ... https://www.frontiersin.org/journals/computational-neuroscience/articles/10.3389/fncom.2025.1556483/pdf
57. Deep learning approaches for EEG-based healthcare ... https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2025.1689073/full
58. A Survey on fMRI-based Brain Decoding ... https://arxiv.org/html/2503.15978v1
59. Multimodal MBC-ATT: cross-modality attentional fusion of ... https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2025.1660532/full
60. Towards AI Search Paradigm https://arxiv.org/html/2506.17188v1
61. Associations between pre-cue parietal alpha oscillations ... https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2025.1625127/full
62. An Enactivist Approach to Human-Computer Interaction https://arxiv.org/pdf/2509.07871
63. Multidimensional EEG features integration with ... https://www.frontiersin.org/journals/psychiatry/articles/10.3389/fpsyt.2025.1624997/full
64. 2023 25K Covering FY2022 | PDF https://www.scribd.com/document/654471576/2023-25K-covering-FY2022
65. glove.6B.100d.txt-vocab.txt https://worksheets.codalab.org/rest/bundles/0xadf98bb30a99476ab56ebff3e462d4fa/contents/blob/glove.6B.100d.txt-vocab.txt
66. LLMs4All: A Review of Large Language Models Across ... https://arxiv.org/html/2509.19580v5
67. Computational approaches to Explainable Artificial ... https://www.sciencedirect.com/science/article/pii/S1566253523002610
68. Review of Machine Learning for Single-Particle Tracking https://pubs.acs.org/doi/10.1021/cbmi.5c00146
69. Mapping the landscape of social behavior https://www.cell.com/cell/fulltext/S0092-8674(25)00154-0?rss=yes
70. Deep Learning in neuroimaging for neurodegenerative ... https://www.sciencedirect.com/science/article/pii/S0022510X25003557
71. Data warehousing methods and processing infrastructure ... https://www.researchgate.net/publication/49699337_Data_warehousing_methods_and_processing_infrastructure_for_brain_recovery_research

72. AACPDM 2024 Free Papers https://onlinelibrary.wiley.com/doi/10.1111/dmcn.16087
73. Mapping EEG Metrics to Human Affective and Cognitive Models https://pmc.ncbi.nlm.nih.gov/articles/PMC12649996/
74. Hierarchical attention enhanced deep learning achieves ... https://www.nature.com/articles/s41598-025-17922-1
75. Advances in non-invasive EEG-based brain-computer ... https://www.sciencedirect.com/science/article/abs/pii/B9780323954372000148
76. (PDF) EEG error-related potentials encode magnitude of ... https://www.researchgate.net/publication/372866183_EEG_error-related_potentials_encode_magnitude_of_errors_and_individual_perceptual_thresholds
77. Brain–Computer Interfaces in Parkinson's Disease ... https://www.mdpi.com/2313-7673/10/8/488
78. Personalized μ-transcranial alternating current stimulation ... https://iopscience.iop.org/article/10.1088/1741-2552/ada980
79. 1 Regulating Next-Generation Implantable Brain-Computer ... https://arxiv.org/pdf/2506.12540
80. 无需手术即可升级脑机接口类似于特斯拉OTA更新 https://www.cls.cn/detail/2261286
81. Brain-computer interface can be upgraded without surgery ... https://www.moomoo.com/news/post/64194629/first-neuralink-participant-reveals-brain-computer-interface-can-be-upgraded
82. First Neuralink participant reveals: Brain-computer ... https://news.futunn.com/en/post/67544300/first-neuralink-participant-reveals-brain-computer-interface-upgrades-can-be
83. Correlates of Theta and Gamma Activity during Visuospatial ... https://pmc.ncbi.nlm.nih.gov/articles/PMC11352628/
84. Theta-Gamma Coupling and Working Memory in ... https://pmc.ncbi.nlm.nih.gov/articles/PMC5911490/
85. Resting-state EEG gamma power predicts immediate and ... https://www.researchgate.net/publication/394977331_Resting-state_EEG_gamma_power_predicts_immediate_and_delayed_recall_in_healthy_adults
86. Out of Rhythm: Compromised Precision of Theta-Gamma ... https://pmc.ncbi.nlm.nih.gov/articles/PMC8896557/
87. EEG Amplitude Modulation Analysis across Mental Tasks https://pmc.ncbi.nlm.nih.gov/articles/PMC10708818/

88. (PDF) Prospects for Implementing Digital Technologies in ... https://www.researchgate.net/publication/387898881_Prospects_for_Implementing_Digital_Technologies_in_the_Administration_of_Justice_in_Zimbabwe
89. Lista de publicaciones del Think Tank del PE https://www.europarl.europa.eu/thinktank/es/research/advanced-search/pdf?keywords=28
90. 机器学习2025_6_6 http://www.arxivdaily.com/thread/68259
91. Timed Automata-Based Strategy for Controlling Drone ... https://www.mdpi.com/2079-9292/13/13/2609
92. Autonomous vehicle control systems - A review of decision ... https://www.researchgate.net/publication/228359982_Autonomous_vehicle_control_systems_-_A_review_of_decision_making
93. 机器学习2025_10_24 http://www.arxivdaily.com/thread/73091
94. Ethical, legal, and policy challenges in field-based ... https://academic.oup.com/jlb/article/11/1/lsae008/7689308
95. (PDF) Free will, neurosciences & robotics https://www.researchgate.net/publication/391865763_Free_will_neurosciences_robotics
96. Arxiv今日论文| 2025-12-24 http://lonepatient.top/2025/12/24/arxiv_papers_2025-12-24
97. The Denario project: Deep knowledge AI agents for ... https://arxiv.org/html/2510.26887v1
98. Who's afraid of synthetic data? Hybrid approaches to ... https://www.sciencedirect.com/science/article/pii/S2352914826000079
99. China https://www.nsfc.gov.cn/Portals/0/fj/fj20230220_01.xlsx
100. 34th Annual Computational Neuroscience Meeting: CNS*2025 https://link.springer.com/article/10.1007/s10827-025-00915-4
101. Spatial Working Memory in Humans Depends on Theta ... https://www.researchgate.net/publication/303633709_Spatial_Working_Memory_in_Humans_Depends_on_Theta_and_High_Gamma_Synchronization_in_the_Prefrontal_Cortex
102. Computer Science https://www.arxiv.org/list/cs/new?skip=25&show=1000
103. Emerging Materials, Wearables, and Diagnostic ... https://www.mdpi.com/2079-6374/12/12/1176
104. Poster Session I | Neuropsychopharmacology https://www.nature.com/articles/npp2013279
105. Poster Abstracts https://journals.lww.com/jisprm/fulltext/2021/04001/poster_abstracts.2.aspx

106. UC San Diego Electronic Theses and Dissertations https://escholarship.org/content/qt7kk2c4nd/qt7kk2c4nd.pdf
107. EEG-informed TMS from manual probing to closed-loop ... https://www.sciencedirect.com/science/article/pii/S1053811925005464
108. Design and implementation of high sampling rate ... https://pmc.ncbi.nlm.nih.gov/articles/PMC10339741/
109. Deep computational neurophenomenology: a methodological ... https://academic.oup.com/nc/article/2025/1/niaf016/8222537
110. Summary of over Fifty Years with Brain-Computer ... https://www.mdpi.com/2076-3425/11/1/43
111. A Unified Frontier in Neuroscience, Artificial Intelligence ... https://arxiv.org/html/2507.10722v1
112. Quantum-like representation of neuronal networks' activity https://www.frontiersin.org/journals/human-neuroscience/articles/10.3389/fnhum.2025.1685339/full
113. 英日略語対訳リスト https://static.aminer.org/pdf/PDF/000/259/939/lara_localization_of_an_automatized_refueling_machine_by_acoustical_sounding.pdf
114. FORMAL METHODS SPECIFICATION AND ... https://www.math.pku.edu.cn/teachers/zhangnx/fm/materials/NASAGB2.pdf
115. A Survey of Papers from Formal Methods for Interactive ... https://enac.hal.science/hal-02364845/document
116. A Survey of Practical Formal Methods for Security https://dl.acm.org/doi/10.1145/3522582
117. Philosophical foundation of the right to mental integrity in ... https://www.researchgate.net/publication/369450299_Philosophical_foundation_of_the_right_to_mental_integrity_in_the_age_of_neurotechnologies
118. MIO-BDT: Construction of Basic Models and Formal ... https://www.mdpi.com/2624-6511/9/1/16
119. On Neurorights - PMC - PubMed Central - NIH https://pmc.ncbi.nlm.nih.gov/articles/PMC8498568/
120. Industrial Strength Formal Using Abstractions https://arxiv.org/pdf/1606.02347
121. Theta-gamma phase-amplitude coupling in auditory cortex ... https://pubmed.ncbi.nlm.nih.gov/36852454/
122. A Survey of Cross-Layer Security for Resource ... https://www.mdpi.com/2076-3417/15/17/9691

123. Brain Computer Interfaces, a Review - PMC - PubMed Central https://pmc.ncbi.nlm.nih.gov/articles/PMC3304110/
124. Critical Review on Brain–Computer Interfaces (BCIs) for ... https://www.researchgate.net/publication/397647166_Critical_Review_on_Brain-Computer_Interfaces_BCIs_for_Effective_Human-Machine_Interactions_HMIs
125. How state-trait continua (body plans, personalities) emerge ... https://www.sciencedirect.com/science/article/pii/S0149763423003718
126. Neural Information Processing https://link.springer.com/content/pdf/10.1007/978-981-96-6576-1.pdf
127. bing.txt - FTP Directory Listing ftp://ftp.cs.princeton.edu/pub/cs226/autocomplete/bing.txt
128. Replay and compositional computation: Neuron https://www.cell.com/neuron/fulltext/S0896-6273(22)01125-4
129. A conditional latent autoregressive recurrent model for ... https://www.nature.com/articles/s41598-024-68944-0
130. Hippocampal sharp wave‐ripple: A cognitive biomarker for ... https://pmc.ncbi.nlm.nih.gov/articles/PMC4648295/
131. Theta‐gamma phase‐amplitude coupling in auditory cortex ... https://pmc.ncbi.nlm.nih.gov/articles/PMC10089097/
132. LNCS 8226 - Neural Information Processing https://link.springer.com/content/pdf/10.1007/978-3-642-42054-2.pdf
133. Beyond neural data: Cognitive biometrics and mental privacy https://www.sciencedirect.com/science/article/pii/S0896627324006524
134. TechDispatch #1/2024 - Neurodata https://www.edps.europa.eu/data-protection/our-work/publications/techdispatch/2024-06-03-techdispatch-12024-neurodata
135. Paradigm Shift in Global Governance of Medical Brain ... https://pmc.ncbi.nlm.nih.gov/articles/PMC12665245/
136. VeriPlan: Integrating Formal Verification and LLMs into End ... https://dl.acm.org/doi/10.1145/3706598.3714113
137. Intersession Robust Hybrid Brain–Computer Interface: Safe ... https://www.mdpi.com/2072-666X/16/11/1264
138. Technologies on the stand: Legal and ethical questions in ... https://www.academia.edu/106998065/Technologies_on_the_stand_Legal_and_ethical_questions_in_neuroscience_and_robotics
139. 脑机芯片可像特斯拉一样OTA 更新，无需二次手术即可获得 ... https://www.ithome.com/0/914/206.htm

140. 无需手术即可升级脑机接口类似于特斯拉OTA更新 https://www.industrysourcing.cn/article/473370
141. Neuralink首例受试者脑机接口已能实现OTA升级，医疗设备 ... https://www.21jingji.com/article/20260121/herald/69ab72b63d0246c2ded20f9165ac4011.html
142. Ramping dynamics and theta oscillations reflect ... https://www.nature.com/articles/s41467-023-44571-7
143. EEGEmotions-27: A Large-Scale EEG Dataset Annotated ... http://ieeexplore.ieee.org/iel8/6287639/10820123/11202184.pdf
144. Investigating the role of oscillations in endogenous and ... https://theses.hal.science/tel-01772802v1/file/2017TOU30026b.pdf
145. Hippocampal sharp wave‐ripple: A cognitive biomarker for ... https://onlinelibrary.wiley.com/doi/full/10.1002%2Fhipo.22488
146. (PDF) ACXNet hybrid deep learning model for cross task ... https://www.researchgate.net/publication/396322908_ACXNet_hybrid_deep_learning_model_for_cross_task_mental_workload_estimation_using_EEG_neural_manifolds
147. Brain-computer interfaces - ISO/IEC JTC 1/SC 43 https://www.iso.org/committee/9082407.html
148. (PDF) Brain-MCP: A Fully Homomorphic Neuro-Symbolic ... https://www.researchgate.net/publication/393129669_Brain-MCP_A_Fully_Homomorphic_Neuro-Symbolic_Protocol_for_Governing_System_3_Cognition_and_Restoring_Semantic_Scarcity
149. ISO/IEC JTC 1/SC 42 - Artificial intelligence https://www.iso.org/committee/6794475.html
150. Poster Session I | Neuropsychopharmacology https://www.nature.com/articles/npp2015325
151. Advancing Neural Networks, Optogenetics, and Cognitive ... https://www.researchgate.net/publication/398781512_Computational_Neuroscience_and_Brain-Computer_Interface_Technologies_Advancing_Neural_Networks_Optogenetics_and_Cognitive_Enhancement_Systems
152. The Israel Society for Neuroscience 19th Annual Meeting Eilat ... https://link.springer.com/article/10.1007/s12031-011-9491-9
153. (PDF) Handbook of Psychophysiology https://www.academia.edu/34699451/Handbook_of_Psychophysiology

154. Brain-computer interfaces - ISO/IEC 8663:2025 https://www.iso.org/standard/83268.html
155. The Theta-Gamma Neural Code - PMC - PubMed Central https://pmc.ncbi.nlm.nih.gov/articles/PMC3648857/
156. The Theta-Gamma Neural Code | Request PDF https://www.researchgate.net/publication/236073864_The_Theta-Gamma_Neural_Code
157. 32nd Annual Computational Neuroscience Meeting: CNS*2023 https://link.springer.com/article/10.1007/s10827-024-00871-5
158. Oral Sessions - 2024 - Journal of Sleep Research https://onlinelibrary.wiley.com/doi/10.1111/jsr.14290
159. Do addictive behaviors explain social inequality regarding ... https://hal.science/hal-04599328v1/file/abstracts_ESBRA2019.pdf
160. Mapping EEG Metrics to Human Affective and Cognitive ... https://www.mdpi.com/2313-7673/10/11/730
161. Internal states as a source of subject-dependent movement ... https://pmc.ncbi.nlm.nih.gov/articles/PMC10687170/
162. The cognitive impacts of large language model interactions ... https://www.frontiersin.org/journals/computational-neuroscience/articles/10.3389/fncom.2025.1556483/full
163. Self-regulation of frontal-midline theta facilitates memory ... https://www.researchgate.net/publication/269166481_Self-regulation_of_frontal-midline_theta_facilitates_memory_updating_and_mental_set_shifting
164. Grounding the neurobiology of language in first principles https://www.sciencedirect.com/science/article/pii/S0010027718301707
165. 26th Annual Computational Neuroscience Meeting (CNS*2017) https://link.springer.com/article/10.1186/s12868-017-0372-1
166. Gravity and the Mind_Alibekov, Gutoreva, Ferre 2025 https://arxiv.org/pdf/2511.05536
167. 26th Annual Computational Neuroscience Meeting (CNS*2017) https://pmc.ncbi.nlm.nih.gov/articles/PMC5592441/
168. Neuroscience Needs Behavior: Correcting a Reductionist ... https://www.researchgate.net/publication/313493106_Neuroscience_Needs_Behavior_Correcting_a_Reductionist_Bias
169. investigating an intensive mindfulness meditation retreat ... https://theses.hal.science/tel-04213136/file/TH2022POUBLAN-COUZARDOTARNAUD.pdf